


# Network Architectures



# LeNet architecture



Tanh as the activation function!



14,197,122 images, 21841 synsets indexed

[Home](#) [Download](#) [Challenges](#) [About](#)


Not logged in. [Login](#) | [Signup](#)

**ImageNet** is an image database organized according to the **WordNet** hierarchy (currently only the nouns), in which each node of the hierarchy is depicted by hundreds and thousands of images. The project has been **instrumental** in advancing computer vision and deep learning research. The data is available for free to researchers for non-commercial use.


**Mar 11 2021. ImageNet website update.**

© 2020 Stanford Vision Lab, Stanford University, Princeton University [imagenet.help.desk@gmail.com](mailto:imagenet.help.desk@gmail.com) Copyright infringement

# AlexNet: The Tipping Point

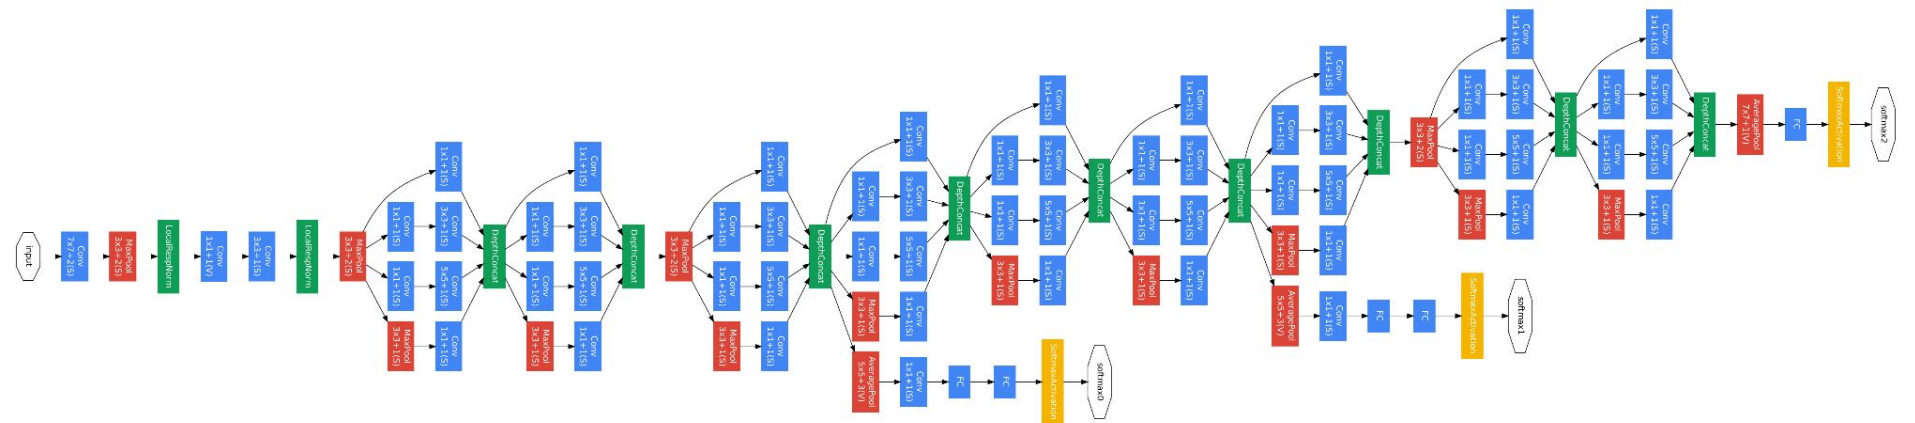


- Convolutional layers with the following kernel sizes:  $11 \times 11$ ,  $5 \times 5$ , and  $3 \times 3$
- Max pooling layers for images downsampling
- Dropout layers to avoid overfitting
- Unlike LeNet, ReLU activation functions in the hidden layers

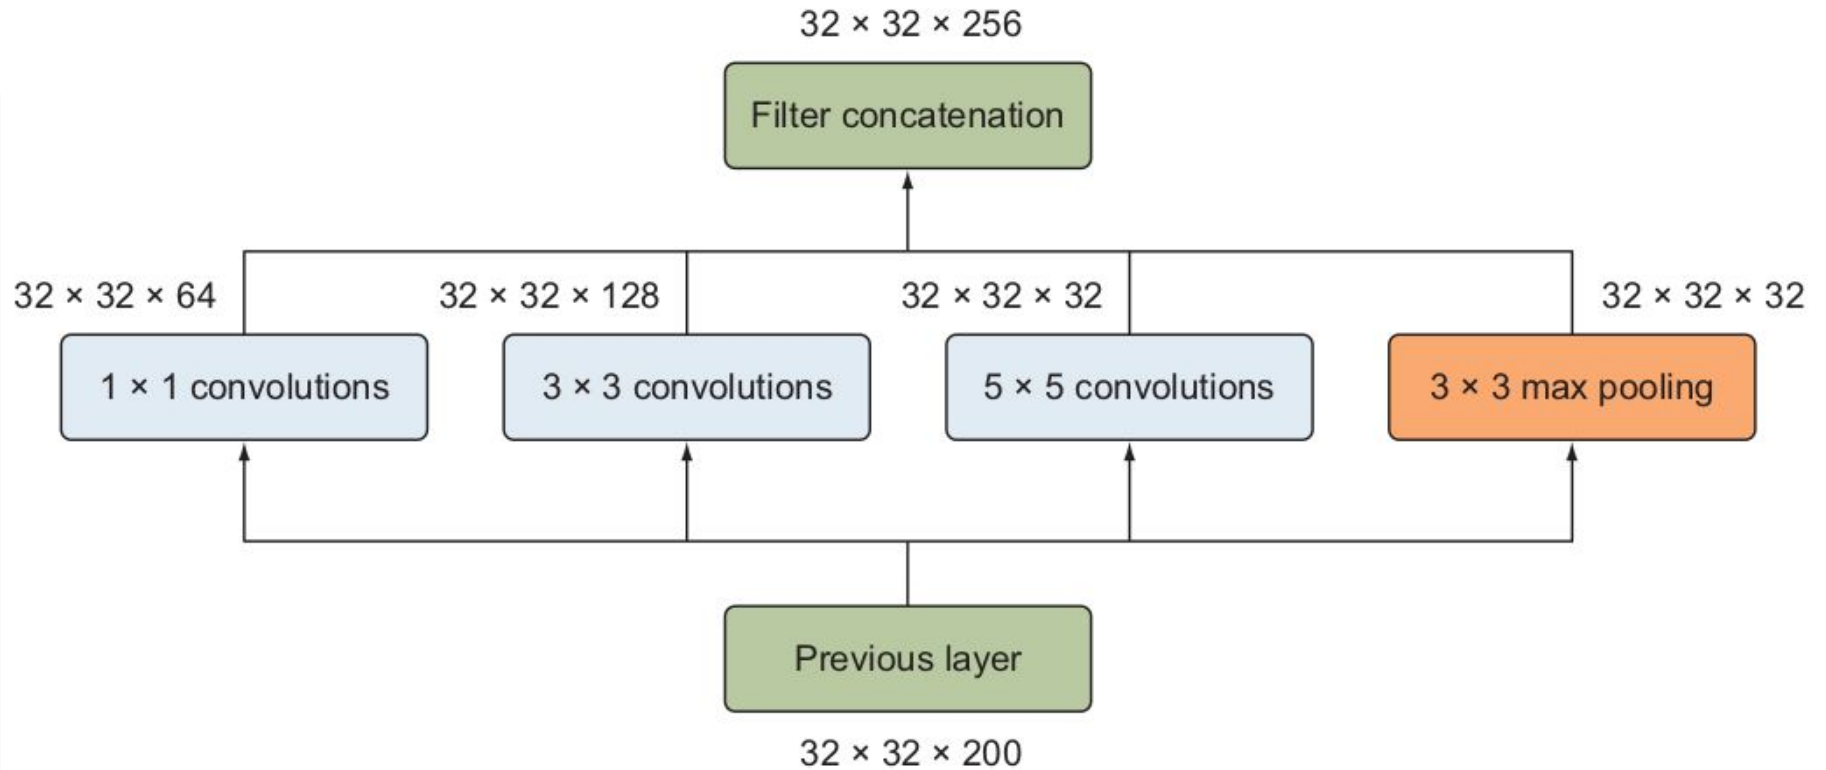




# Inception and GoogLeNet




## Inception Module: Naive Approach






## $1 \times 1$ Convolutional (Bottleneck) Layers




$6 \times 6 \times 32$

\*



$1 \times 1 \times \# \text{ filters}$

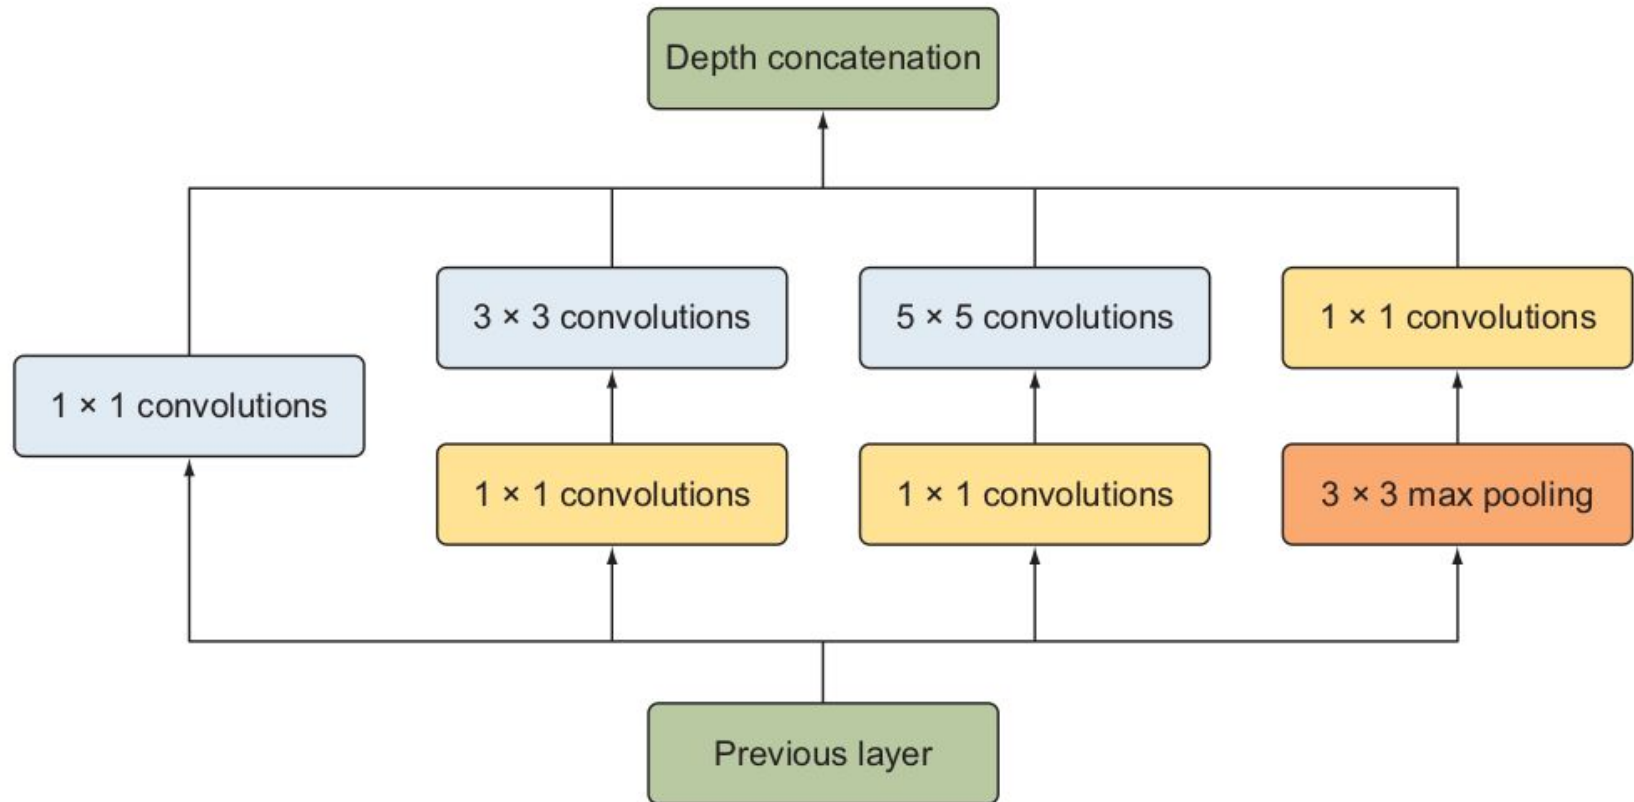
=



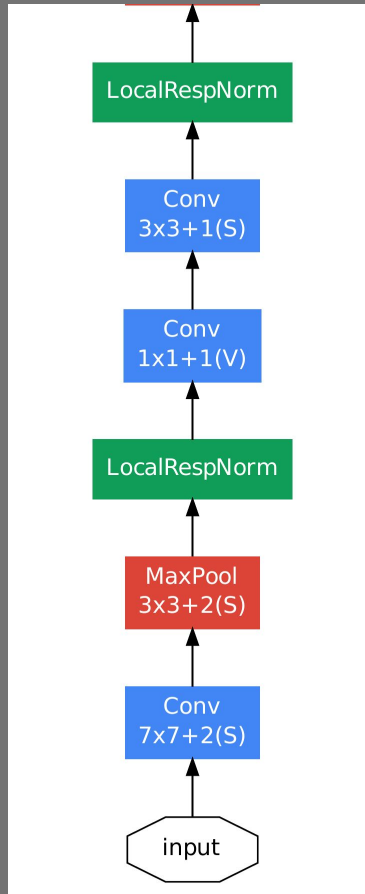
$6 \times 6 \times \# \text{ filters}$

**$1 \times 1$  conv layers preserve the spatial dimensions but change the depth.**

# Inception Module




# Inception and GoogLeNet



# ResNet


VGG-19



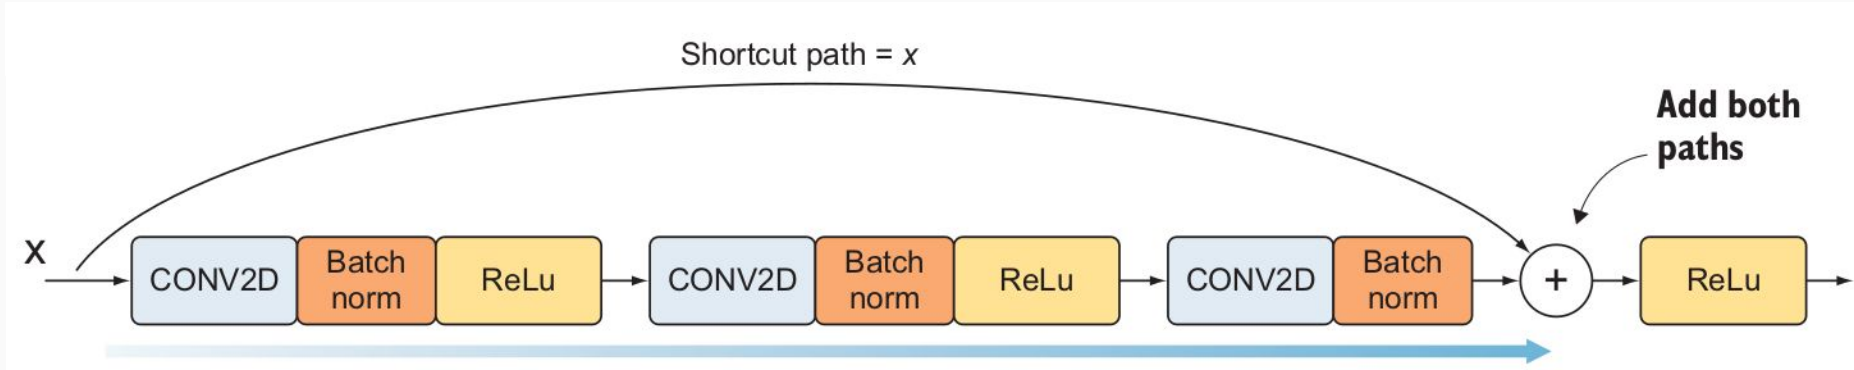
34-layer plain



34-layer residual



## Residual Block: Same input and output



# Residual Block

