

# Introduction to Convolutional Neural Networks (CNNs)

Farhad Maleki

# Agenda

- Introduction to Convolutional Neural Networks
- Convolution Operation
- Activation Functions
- Pooling Layers
- Batch Normalization
- Regularization Techniques
- Transfer Learning
- Advanced Architectures
- Real-world Applications

# Introduction to Convolutional Neural Networks

- CNNs are designed to process grid-like data, such as images
- They exploit the **spatial structure** of images using **local connectivity** and **parameter sharing**
- CNNs are capable of **hierarchical feature learning**

Why not Using MLP for Images?

# Why CNNs are Commonly Used in Computer Vision


- Scalability: High dimension of visual data
- Local connectivity
- Parameter sharing
- Hierarchical feature learning
- Translation invariance



# Convolution Operation

- Convolution involves sliding a filter over the input image
- Computes the element-wise product between the filter and the image patch
- Outputs a feature map or activation map


# Convolution Operation



[Link to the Visualizations](#)

```
nn.Conv2d(in_channels=1, out_channels=1, kernel_size=3)
```


## Convolution with Padding



```
nn.Conv2d(in_channels=1, out_channels=1, kernel_size=4, padding=2)
```




## Convolution Half Padding (Same Padding)




```
nn.Conv2d(in_channels=1, out_channels=1, kernel_size=3, padding=1)
```

## Convolution with Stride-No Padding




```
nn.Conv2d(in_channels=1, out_channels=1, kernel_size=3, stride=2, padding=0)
```

## Convolution with Stride and Padding




```
nn.Conv2d(in_channels=1, out_channels=1, kernel_size=3, stride=2, padding=1)
```

# Dilated Convolution



```
nn.Conv2d(in_channels=1, out_channels=1, kernel_size=3, stride=1, padding=0, dilation=2)
```

# Transposed convolution



```
nn.ConvTranspose2d(in_channels=1, out_channels=1, kernel_size=3, stride=1, padding=2, dilation=2)
```

# Activation Functions

- Introduce non-linearity into the model
- Common activation functions: ReLU, sigmoid, and tanh
- ReLU is computationally efficient and helps mitigate the vanishing gradient problem

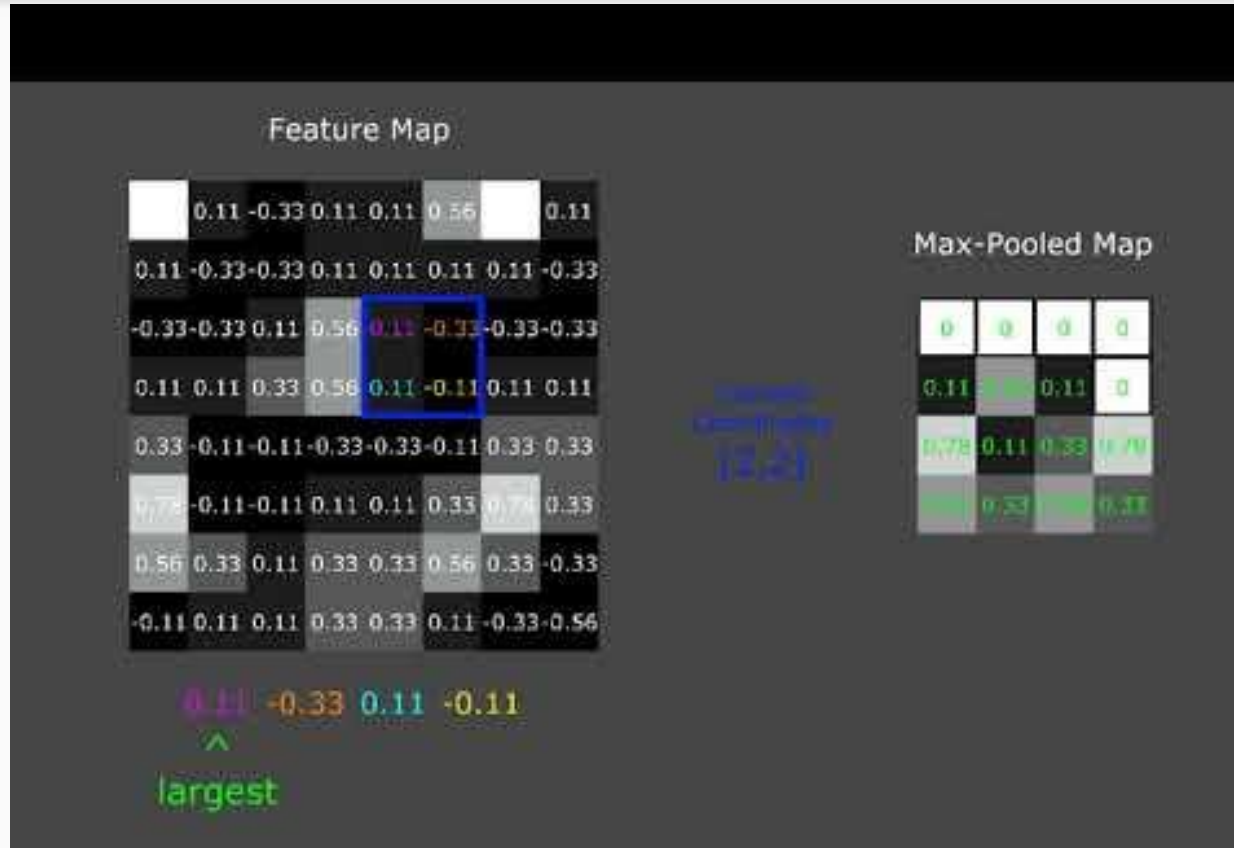
# Pooling Layer

- To reduce the spatial dimensions of the feature maps
- To reduce the number of parameters in the network and control overfitting
- Used to downsample spatial dimensions and reduce computational complexity
- Increases translation invariance
- Common types: max pooling and average pooling

`nn.MaxPool2d(kernel_size)`

`nn.AvgPool2d(kernel_size)`

# Max-Pooling



[Video Link](#)

`nn.MaxPool2d(kernel_size=2)`



# The covariate shift problem


Input layer

L1

L2

L3

L4



# Batch normalization

$$\mu_B \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad \longleftarrow \text{Mini-batch mean}$$

$$\sigma_B^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_B)^2 \quad \longleftarrow \text{Mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta$$

# Batch Normalization

- Improves training by normalizing input distribution of each layer
- Mitigates the internal covariate shift problem
- Improves training speed, generalization, and stability

`nn.BatchNorm2d(num_features)`